

Quantum Computing, Quantum Machine Learning and Quantum Information Theories

Morten Hjorth-Jensen¹

Department of Physics and Center for Computing in Science Education,
University of Oslo, Norway¹

April 30

© 1999-2025, Morten Hjorth-Jensen. Released under CC Attribution-NonCommercial 4.0 license

Plans for the week of April 28-May 2

1. Classical Support Vector Machines, reminder from last week
2. Classical Kernels and transition to Quantum Kernels
3. Quantum Support Vector Machines
4. And if we get time, we may start with Quantum Neural Networks

Reading recommendations

1. For classical Support Vector Machines, see lecture notes by Morten Hjorth-Jensen at https://compphysics.github.io/MachineLearning/doc/LectureNotes/_build/html/chapter5.html
2. Claudio Conti, Quantum Machine Learning (Springer), sections 1.5-1.12 and chapter 2, see <https://link.springer.com/book/10.1007/978-3-031-44226-1>.
3. See also Maria Schuld's article at <https://arxiv.org/abs/2101.11020>.
4. For coding, we included examples using PennyLane at <https://pennylane.ai/> and Qiskit, see <https://qiskit-community.github.io/qiskit-machine-learning/>

What is Quantum Machine Learning?

Quantum Machine Learning (QML) integrates quantum computing with machine learning algorithms to exploit quantum advantages. It explores how quantum computing can enhance classical machine learning.

Motivation:

1. High-dimensional Hilbert spaces for better feature representation.
2. Quantum parallelism for faster computation.
3. Quantum entanglement for richer data encoding.

Quantum Speedups in ML

Why Quantum?

1. **Quantum Parallelism:** Process multiple states simultaneously.
2. **Quantum Entanglement:** Correlated states for richer information.
3. **Quantum Interference:** Constructive and destructive interference to enhance solutions.

Challenges in Quantum Machine Learning

Quantum Hardware Limitations:

1. Noisy Intermediate-Scale Quantum (NISQ) devices.
2. Decoherence and limited qubit coherence times.

Data Encoding:

1. Efficient embedding of classical data into quantum states.

Scalability:

1. Difficult to scale circuits to large datasets.

Support vector machines

As discussed last week, a central model in classical supervised learning is the support vector machine (SVM), which is a max-margin classifier. SVMs are widely used for binary classification and have extensions to regression problems as well. They build on statistical learning theory and are known for finding decision boundaries with maximal margin. In particular, SVMs can perform non-linear classification by employing the kernel trick, which implicitly maps data into a high-dimensional feature space via a kernel function.

The quantum SVM

A Quantum Support Vector Machine (QSVM) replaces the classical kernel or feature map with a quantum procedure. In QSVM, classical data points $\mathbf{x}$ are encoded into quantum states $|\phi(\mathbf{x})\rangle$ via a quantum feature map (a parameterized quantum circuit). The inner product (overlap) between two such states serves as a quantum kernel, measuring data similarity in a high-dimensional Hilbert space. Many quantum models, including **quantum neural networks**, ultimately behave like kernel methods: they analyze data in exponentially-large Hilbert spaces, accessible only through state overlaps, see for example Maria Schuld, Supervised quantum machine learning models are kernel methods, at [arXiv:2101.11020](https://arxiv.org/abs/2101.11020)

Schuld shows that supervised quantum models can be reformulated as SVMs whose kernels compute the distance (inner product) between data-encoded quantum states. Thus, QSVMs can leverage the expressive power of quantum state space to potentially separate data that classical kernels struggle with.

Classical Support Vector Machines (SVMs)

Support Vector Machines (SVM) are supervised learning algorithms used for classification tasks, and regression tasks as well. The main goal of SVMs is to find the optimal separating hyperplane (in high-dimensional space) that provides a maximum margin between classes.

First mathematical formulation of SVMs (formal details below)

For a dataset $(\mathbf{x}_i, y_i)$ where $\mathbf{x}_i \in \mathbb{R}^n$ and $y_i \in \{-1, 1\}$, the decision boundary is defined as:

$$f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b = 0.$$

The goal is to optimize

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2,$$

subject to

$$y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1, \quad \forall i$$

Kernels and more

In classical SVMs, kernels help with non-linear data separations. Quantum computers can speed up the computation of complex kernel evaluations by efficiently simulating an inner product in an exponentially large Hilbert space.

Quantum Support Vector Machines (QSVMs)

Quantum Kernel Estimation:

1. Maps classical data to a quantum Hilbert space.
2. Quantum kernel measures similarity in high-dimensional space.

Quantum-enhanced kernel:

$$K(\mathbf{x}, \mathbf{x}') = |\langle \psi(\mathbf{x}) | \psi(\mathbf{x}') \rangle|^2$$

Here, $|\phi(\mathbf{x})\rangle$ is the quantum state encoding of the classical data $\mathbf{x}$.
Advantage: Potentially exponential speedup over classical SVMs.

Quantum Neural Networks (QNNs)

Quantum Neural Networks replace classical neurons with parameterized quantum circuits.

Key Concepts:

1. Quantum Gates as Activation Functions.
2. Variational Quantum Circuits (VQCs) for optimization.

Parameterized Quantum Circuit:

$$U(\theta) = \prod_i R_y(\theta_i) \cdot CNOT \cdot R_x(\theta_i)$$

Advantage: Quantum gradients enable exploration of non-convex landscapes.

Quantum Boltzmann Machines (QBMs)

Quantum Boltzmann Machines leverage quantum mechanics to sample from a probability distribution.

1. Quantum tunneling aids in escaping local minima.
2. Quantum annealing for optimization problems.

Quantum Hamiltonian:

$$H = - \sum_i b_i \sigma_i^z - \sum_{ij} w_{ij} \sigma_i^z \sigma_j^z$$

Advantage: Efficient sampling in complex probability distributions.

Quantum SVMs

The idea of a classical SVM is that we have a set of points that are in either one group or another and we want to find a line that separates these two groups. This line can be linear, but it can also be much more complex, which can be achieved by the use of Kernels.

What will we need in the case of a quantum computer?

We will have to translate the classical data point $\vec{x}$ into a quantum datapoint $|\Phi(\vec{x})\rangle$. This can be achieved by a circuit $\mathcal{U}_{\Phi(\vec{x})}|0\rangle$. Here $\Phi()$ could be any classical function applied on the classical data $\vec{x}$.

We need a parameterized quantum circuit $W(\theta)$ that processes the data in a way that in the end we can apply a measurement that returns a classical value -1 or 1 for each classical input $\vec{x}$ that identifies the label of the classical data.

The most general ansatz

Following these steps we can define an ansatz for this kind of problem which is

$$W(\theta)\mathcal{U}_{\Phi}(\vec{x})|0\rangle.$$

These kind of ansatzes are called quantum variational circuits.

Quantum SVM

In the case of a quantum SVM we will only use the quantum feature maps

$$\mathcal{U}_{\Phi(\vec{x})},$$

to translate the classical data into quantum states and build the Kernel of the SVM out of these quantum states. After calculating the Kernel matrix on the quantum computer we can train the Quantum SVM the same way as the classical SVM.

Defining the Quantum Kernel

The idea of the quantum kernel is exactly the same as in the classical case. We take the inner product

$$K(\vec{x}, \vec{z}) = |\langle \Phi(\vec{x}) | \Phi(\vec{z}) \rangle|^2 = \langle 0^n | \mathcal{U}_{\Phi(\vec{x})}^\dagger \mathcal{U}_{\Phi(\vec{z})} | 0^n \rangle,$$

but now with the quantum feature maps

$$\mathcal{U}_{\Phi(\vec{x})}.$$

The idea is that if we choose a quantum feature maps that is not easy to simulate with a classical computer we might obtain a quantum advantage.

Side note

There is no proof yet that the QSVM brings a quantum advantage, but the argument the authors of <https://arxiv.org/pdf/1804.11326.pdf> make, is that there is for sure no advantage if we use feature maps that are easy to simulate classically, because then we would not need a quantum computer to construct the Kernel.

Feature map

For the feature maps we use the ansatz

$$\mathcal{U}_{\Phi(x)} = U_{\Phi(x)} \otimes H^{\otimes n},$$

where

$$U_{\Phi(x)} = \exp \left(i \sum_{S \in n} \phi_S(x) \prod_{i \in S} Z_i \right),$$

which simplifies a lot when we (like in

<https://arxiv.org/pdf/1804.11326.pdf>) only consider $S \leq 2$ interactions, which means we only let two qubits interact at a time. For $S \leq 2$ the product $\prod_{i \in S}$ only leads to interactions $Z_i Z_j$ and non interacting terms Z_i . And the sum $\sum_{S \in n}$ over all these terms that are possible with n qubits.

Classical functions

Finally we define the classical functions $\phi_i(\vec{x}) = x_i$ and $\phi_{i,j}(\vec{x}) = (\pi - x_i)(\pi - x_j)$.

If we write this ansatz for 2 qubits and $S \leq 2$ we see how it simplifies:

$$U_{\Phi(x)} = \exp(i(x_1 Z_1 + x_2 Z_2 + (\pi - x_1)(\pi - x_2) Z_1 Z_2)).$$

We won't get into details too much here, why we would take this ansatz. It is simply an ansatz that is simple enough and leads to good results.

Finally we can define a depth of these circuits. Depth 2 means we repeat this ansatz two times. Which means our feature map becomes $U_{\Phi(x)} \otimes H^{\otimes n} \otimes U_{\Phi(x)} \otimes H^{\otimes n}$.

Plans for next week

1. Quantum neural networks

Key properties of SVMs

Support vector machines use only the support vectors (training points on the margin) to define the classifier, making the decision function sparse and often memory-efficient. They are robust to overfitting when properly regularized (max-margin helps generalization) .

However, training SVMs can be costly for very large datasets, as one must solve a quadratic program that scales roughly as $O(N^3)$ in naive implementations. Evaluating the kernel (especially for complex kernels or large feature dimension) can also be expensive. In fact, “kernel methods for machine learning are ubiquitous for pattern recognition, with SVMs being the most well-known method...” , but they face limitations when feature spaces become large and kernel computations costly. This sets the stage for using quantum computers to estimate kernels or implement SVMs more efficiently.

Quantum Kernels and Feature Maps

A core idea of QSVM is to use a quantum feature map that encodes classical input $\mathbf{x}$ into a quantum state $|\phi(\mathbf{x})\rangle$ in an exponentially large Hilbert space.

Mathematically, a quantum feature map is a unitary circuit $U(\mathbf{x})$ acting on an n -qubit register initialized to $|0\rangle^{\otimes n}$, producing

$$|\phi(\mathbf{x})\rangle = U(\mathbf{x})|0\rangle^{\otimes n}.$$

Input dependence

This state depends on the input $\mathbf{x} \in \mathbb{R}^d$. The choice of $U(\mathbf{x})$ is analogous to choosing a classical feature mapping $\phi(\mathbf{x})$ in kernel methods. Common quantum feature maps include amplitude encoding (embedding data amplitudes in the state vector) and angle encoding (rotations whose angles are functions of $\mathbf{x}$) . Once data are encoded as quantum states, one defines a quantum kernel as an overlap between states. A simple definition is the fidelity:

$$k(\mathbf{x}, \mathbf{x}') = |\langle \phi(\mathbf{x}) | \phi(\mathbf{x}') \rangle|^2.$$

Quantum kernels

That is, the kernel is the squared magnitude of the inner product of the two quantum states. Another common (unnormalized) version is

$$k'(\mathbf{x}, \mathbf{x}') = \langle \phi(\mathbf{x}) | \phi(\mathbf{x}') \rangle,$$

but measuring this amplitude directly can be difficult due to the absolute value. Many QSVM implementations (and theoretical definitions) use the squared overlap. In any case, the kernel measures similarity: if $|\phi(\mathbf{x})\rangle$ and $|\phi(\mathbf{x}')\rangle$ are close in Hilbert space, $k(\mathbf{x}, \mathbf{x}')$ is large.

What is a quantum kernel?

A quantum kernel is a function that measures the similarity between two quantum states, which are obtained by applying a quantum feature map to classical data. In other words, each classical vector $\mathbf{x}$ is mapped to a quantum state $|\phi(\mathbf{x})\rangle$ via a feature map, and the kernel is defined by the inner product of these states. Concretely, one may write

$$K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j) = |\langle \phi(\mathbf{x}_i) | \phi(\mathbf{x}_j) \rangle|^2.$$

This forms a positive semidefinite kernel matrix K on the dataset, just as in classical SVMs.

Power of quantum kernels

The power of quantum kernels comes from the exponentially large feature space. An n -qubit state has dimension 2^n , so even a simple feature map using n qubits corresponds to an embedding of $\mathbf{x}$ into $\mathbb{C}^{2^n}$. Entangling gates in $U(\mathbf{x})$ allow highly non-linear features. In principle, a carefully chosen quantum feature map can capture complex structures in data that are hard to model classically. For example, Havlíček et al. used a feature map involving Pauli-Z rotations and controlled-Z gates to embed data into a 2^n -dimensional space. They note that classical kernels become expensive as feature dimensions grow, while quantum circuits naturally operate in these large spaces.

Estimating quantum kernels

We must be able to estimate the quantum kernel in practice. Given two data vectors $\mathbf{x}$ and $\mathbf{x}'$, we need to compute (or measure) $|\langle\phi(\mathbf{x})|\phi(\mathbf{x}')\rangle|^2$. This can be done on a quantum computer by preparing $|\phi(\mathbf{x})\rangle$ and $|\phi(\mathbf{x}')\rangle$ in some interference experiment. For instance, one can prepare $|\phi(\mathbf{x})\rangle$, then apply $U(\mathbf{x}')^\dagger$, and measure the probability of returning to the $|0\rangle^{\otimes n}$ state. That probability is exactly $|\langle 0|U(\mathbf{x})^\dagger U(\mathbf{x}')|0\rangle|^2 = |\langle\phi(\mathbf{x})|\phi(\mathbf{x}')\rangle|^2$. In PennyLane, this is conveniently done by defining a QNode circuit that encodes $\mathbf{x}$, then encodes $\mathbf{x}'$ inversely (via the adjoint).

Code example

For example, using PennyLane's AngleEmbedding template, we can write:

```
import pennylane as qml
from pennylane import numpy as np
dev = qml.device('default.qubit', wires=2)
@qml.qnode(dev)
def kernel_circuit(x1, x2):
    qml.templates.AngleEmbedding(x1, wires=[0,1])
    qml.adjoint(qml.templates.AngleEmbedding)(x2, wires=[0,1])
    return qml.probs(wires=[0,1])[0]
```

What happens

Here, **AngleEmbedding(x1)** encodes the vector **x1** into two qubits via rotations, and `qml.adjoint` applies the inverse embedding for **x2**. The output probability of measuring the $|00\rangle$ state is exactly the squared overlap of the two embedded states. Using this kernel function, PennyLane can compute the kernel matrix $K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$ efficiently for training. In fact, PennyLane provides a utility `qml.kernels.kernel_matrix` to evaluate such kernels systematically .

Quantum feature map

A quantum feature map $\mathbf{x} \mapsto |\phi(\mathbf{x})\rangle$ combined with the kernel $k(\mathbf{x}, \mathbf{x}') = |\langle \phi(\mathbf{x}) | \phi(\mathbf{x}') \rangle|^2$ defines a QSVM kernel. The intuition is that the quantum device is implicitly computing a rich similarity measure via its high-dimensional state space.

This is analogous to the classical kernel trick but potentially more powerful: **quantum kernels can potentially offer advantages over classical kernel methods, such as speedup, accuracy, flexibility, and robustness.**

Examples of Feature Map Circuits

There are many choices for $U(\mathbf{x})$. Common simple maps include:

Angle (or rotation) embedding:

For an n -dimensional feature vector $\mathbf{x} = (x_1, \dots, x_n)$, apply single-qubit rotations $R_X(x_i)$ or $R_Y(x_i)$ to the i th qubit. For example:

$$U(\mathbf{x}) = \bigotimes_i R_Y(x_i) = R_Y(x_1) \otimes \cdots \otimes R_Y(x_n).$$

This maps x_i to an angle on the Bloch sphere. It produces product states (no entanglement).

More features

Amplitude encoding:

Encode the normalized data vector $\mathbf{x}$ directly into the amplitudes of the quantum state. For example, for 2 qubits one can map (x_1, x_2, x_3, x_4) (with $\sum x_i^2 = 1$) to $x_1|00\rangle + x_2|01\rangle + x_3|10\rangle + x_4|11\rangle$. This uses exponentially many amplitudes and can be powerful, but requires complex circuits.

Entangling feature maps:

Combine rotations with entangling gates. For instance, the ZZ feature map used in [Havlíček et al.] alternates single-qubit $R_Z(x_i)$ rotations with controlled-Z (CZ) gates between qubits. This creates entangled states whose overlaps can capture correlations in the data. In PennyLane, one could implement something like:

```
for i in range(n):
   qml.RZ(x[i], wires=i)
for i in range(n-1):
   qml.CNOT(wires=[i, i+1])
for i in range(n):
   qml.RZ(x[i], wires=i)
```

Polynomial feature map

These are inspired by classical polynomial kernels, one can encode quadratic or higher terms. For example, apply $R_Z(x_i)$, $R_X(x_i)$ on each qubit and use multi-qubit rotations to create cross-terms . The choice of feature map affects the kernel. A richer (more entanglement) map may separate classes better but might also be harder to learn or require more quantum resources. In practice, one often starts with simple maps (e.g. AngleEmbedding) and experiments.

Kernel Estimation on Quantum Devices

Once $|\phi(\mathbf{x})\rangle$ is defined, we must compute $k(\mathbf{x}, \mathbf{x}')$ on actual quantum hardware. Typically this involves running a quantum circuit multiple times (shots) to estimate a probability. For the overlap kernel, one measures the probability of a certain outcome after applying $U(\mathbf{x})$ and $U(\mathbf{x}')^\dagger$. Alternative schemes include the SWAP test, which uses an ancilla qubit to directly estimate overlaps, but this requires extra gates. The above adjoint-circuit method only needs the same number of qubits as the feature map.

Short summary

In summary, the quantum kernel is given by inner products in a quantum-defined feature space. We emphasize: all computations on the quantum device yield kernel values (scalars). The heavy lifting of classification (optimizing α_i) remains a classical task, typically done by classical SVM software using the quantum-computed kernel matrix. This **hybrid** approach leverages quantum hardware where it's needed (kernel evaluation) and classical hardware for optimization.

Quantum Kernel Formula: Starting from a feature map circuit $U(\mathbf{x})$, show explicitly how to compute $k(\mathbf{x}, \mathbf{x}') = |\langle \phi(\mathbf{x}) | \phi(\mathbf{x}') \rangle|^2$ by a quantum circuit and measurements. (Hint: prepare $|\phi(\mathbf{x})\rangle$, apply $U(\mathbf{x}')^\dagger$, measure probability of $|0 \cdots 0\rangle$.)

Properties of Quantum Kernels: Prove that the quantum kernel $K(\mathbf{x}, \mathbf{x}') = |\langle \phi(\mathbf{x}) | \phi(\mathbf{x}') \rangle|^2$ is positive semidefinite (i.e. a valid Mercer kernel). Example

Calculation: Consider a single-qubit angle embedding

$|\phi(x)\rangle = \cos(x/2)|0\rangle + \sin(x/2)|1\rangle$. Compute the kernel

$k(x, x') = |\langle \phi(x) | \phi(x') \rangle|^2$ in closed form. What function of x, x' do you get? Entangled Feature Map: Suppose $n = 2$ qubits and define

$U(\mathbf{x}) = R_{\mathbf{X}}(x_1) \otimes R_{\mathbf{X}}(x_2)$ (no entanglement). What classical kernel

Quantum Support Vector Machine Theory

We now combine the classical SVM formulation with quantum kernels. In essence, we run a classical SVM algorithm (e.g. solving the dual) but supply it with a quantum-computed kernel matrix.

The resulting model is the same form:

$f(\mathbf{x}) = \text{sign}(\sum_i \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b)$, but K comes from a quantum device. The hope is that the quantum kernel captures data structure that a classical kernel might not.

QSVM Formulation via Quantum Kernels

Given training data $(\mathbf{x}_i, y_i)$, we choose a quantum feature map $U(\mathbf{x})$ and define the kernel $K_{ij} = |\langle \phi(\mathbf{x}_i) | \phi(\mathbf{x}_j) \rangle|^2$. Then we solve the binary type of SVM problems discussed last week, with $\mathbf{x}_i^T \mathbf{x}_j$ replaced by K_{ij} . That is:

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j),$$

subject to $\sum_i \alpha_i y_i = 0$, $0 \leq \alpha_i \leq C$. After solving for α_i , the decision function is

$$f(\mathbf{x}) = \text{sign}\left(\sum_i \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b\right).$$

All kernel values $K(\mathbf{x}_i, \mathbf{x}_j)$ are estimated by the quantum device. Thus the QSVM is essentially a classical SVM with a learned quantum kernel.

Important training point

A key point is that training is still classical: one uses a conventional convex solver (e.g. libSVM) on the kernel matrix. The quantum part is in computing K_{ij} . This is sometimes called a quantum kernel estimator. Alternatively, one can approximate gradients of a parameterized kernel for optimization, but here we focus on the fixed-kernel case.

Why might a quantum kernel be better than a classical one?

The hope is that the Hilbert-space embedding is very high-dimensional (exponential), so that more functions (decision boundaries) become linear in that space. Replacing dot products by quantum overlaps is a generalization of classical kernels. In fact, Schuld argues that many near-term **quantum neural networks** are really kernel machines: “a lot of near-term and fault-tolerant quantum models can be replaced by a general support vector machine whose kernel computes distances between data-encoding quantum states” . Thus, one may as well solve directly in kernel form.

Quantum SVM Algorithms (large-scale vs NISQ)

There are two broad approaches in the literature for QSVMs:

Quantum Algorithms (fault-tolerant):

Early work by Rebentrost, Mohseni, and Lloyd (2014) formulated an SVM in terms of quantum linear algebra. They showed that one can invert the kernel matrix (a positive semidefinite matrix) using quantum algorithms (HHL algorithm) in time polylogarithmic in N and d . Concretely, they assumed quantum RAM (QRAM) access to data and used a quantum subroutine to solve the dual SVM as a linear system, yielding the vector of α_i in superposition. Under ideal conditions (sparse data, good condition number), this yields exponential speed-up over classical SVM training. However, it requires fully fault-tolerant quantum hardware and QRAM, making it more a theoretical result than a near-term implementation.

And NISQ Quantum Kernels

More recent work focuses on near-term devices. Here the idea is to compute the kernel matrix entries with a quantum circuit on current hardware or simulators, and then use a classical SVM solver for the rest. This does not claim an exponential runtime speed-up (indeed, the kernel matrix still has N^2 entries). Instead, it aims for a possible quality or feature-based advantage: the quantum feature map might separate data better than any known classical kernel. This approach has been demonstrated on small datasets (e.g. Iris) and studied theoretically. For example, Havlicek et al. showed on a toy problem that a quantum kernel can correctly classify points that a simple classical kernel cannot. However, other studies have found that for random data classical kernels often suffice, so the advantage is subtle. Research is ongoing on which problems truly benefit from quantum kernels.

Quantum neural network

Another variation is the quantum variational classifier, sometimes called a quantum neural network. Instead of precomputing a fixed kernel, one trains a parameterized quantum circuit to output labels. Interestingly, Schuld (2021) shows that variational quantum models, when trained by minimizing a loss, are mathematically equivalent to kernel machines with a particular kernel determined by the circuit. In fact, one can often find a kernel SVM that matches or outperforms the variational model. In practice, one can combine these: use a trainable quantum embedding $U(\mathbf{x}; \theta)$ with tunable parameters θ , and optimize θ to maximize the SVM classification accuracy. This is called a quantum kernel learning approach.

Quantum vs Classical SVM Performance

For our purposes, we concentrate on the case where the quantum feature map is fixed (or manually chosen), and we use classical SVM optimization on the resulting kernel.

Quantum SVMs are subject to the same generalization considerations as classical SVMs: overfitting can occur if the kernel is too flexible, and the choice of regularization C matters. One advantage is that quantum kernels can in principle provide very complex decision boundaries. On the other hand, estimating the kernel matrix accurately requires many quantum measurements (shots). Each kernel entry K_{ij} is obtained by sampling; statistical noise in kernel values can degrade SVM performance. Error mitigation techniques may be needed. Also, evaluating a kernel classically might in some cases simulate the quantum one if the circuit is not too complex; then no advantage is gained.

Comparing computational complexity

1. Classical SVM: Training roughly $O(N^3)$ time, inference $O(N_s)$ time per point, where N_s is number of support vectors.
2. Quantum kernel SVM: Kernel estimation takes $O(n_q N^2)$ time (where n_q is quantum circuit time per inner product) to compute all K_{ij} , plus classical $O(N^3)$. No asymptotic speed-up unless one uses quantum linear algebra (which still needs QRAM).
3. Quantum linear SVM: Potential $O(\log N)$ for training under ideal QRAM assumptions, inference also $O(\log N)$, at expense of large constant overhead.

Actual implementations

In practice on current hardware, QSVM speed is dominated by circuit execution time. However, QSVM experiments are valuable to test expressivity: for some data, a simple quantum feature map yields perfect classification where classical kernels fail. For example, the **Swiss roll** or concentric circle datasets are often used as quantum kernel benchmarks.

Finally, it is worth noting that both the variational quantum classifier and the kernel QSVM ultimately hinge on how data is encoded into quantum states. Schuld summarizes that **the way that data is encoded into quantum states is the main ingredient that can potentially set quantum models apart from classical machine learning models**. Thus, a lot of focus is on designing powerful quantum embeddings and kernels. The SVM framework provides a clear way to evaluate them.

Practical Implementation with PennyLane

We now turn to hands-on implementation of QSVM using PennyLane, a Python library for quantum machine learning. PennyLane interfaces with quantum simulators and hardware and provides convenient utilities for kernels and optimization. In this chapter we will outline how to construct a QSVM: define a quantum feature map, compute the kernel matrix, and train a classical SVM using scikit-learn.

Setting Up PennyLane

First, install PennyLane (pip install pennylane) and import necessary modules:

```
import pennylane as qml
from pennylane import numpy as np
from sklearn import svm
```

PennyLane uses devices to run quantum circuits. For QSVM experiments we can use the default.qubit simulator:

```
dev = qml.device("default.qubit", wires=n_qubits, shots=None)
```

Here `n_qubits` is chosen based on our feature map (often equal to data dimension or its log, etc.) `shots=None` means analytic probabilities (no sampling noise). For realistic experiments, set `shots=100` or so to simulate finite sampling.

Defining Quantum Feature Maps in PennyLane

We must define a quantum circuit (QNode) that encodes input vectors. For example, for a 2-dimensional input $\mathbf{x} = (x_0, x_1)$, we might use AngleEmbedding:

```
@qml.qnode(dev)
def feature_map(x):
    qml.templates.AngleEmbedding(x, wires=[0,1])
    return [qml.expval(qml.PauliZ(0)), qml.expval(qml.PauliZ(1))]
```

This feature map returns expectation values (though for kernel computation we will use a different approach). Alternatively, we can directly define a unitary:

```
@qml.qnode(dev)
def U(x):
    qml.RY(x[0], wires=0)
    qml.RY(x[1], wires=1)
    return qml.state()
```

This QNode returns the quantum state (though `qml.state()` may require a special device). For kernel computation, we actually want to overlap two circuits, as shown below.

Computing Quantum Kernel Matrices

To train an SVM, we need the kernel matrix $K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$.

PennyLane provides `qml.kernels.kernel_matrix`, which takes two datasets and a kernel function. We must supply a function `kernel(x1,x2)` that returns the overlap of states.

Using the adjoint trick, we define:

```
@qml.qnode(dev)
def kernel_circuit(x1, x2):
    # encode x1
    qml.templates.AngleEmbedding(x1, wires=[0,1])
    # apply inverse embedding of x2
    qml.adjoint(qml.templates.AngleEmbedding)(x2, wires=[0,1])
    # return probability of |00>
    return qml.probs(wires=[0,1])[0]
```

This returns $|\langle \phi(x_1) | \phi(x_2) \rangle|^2$ (the prob. of measuring all zeros).

We then set:

```
kernel = lambda a, b: kernel_circuit(a, b)
```

and compute kernel matrices. For training:

```
X_train = np.array([...]) # shape (N_train, 2)
Y_train = np.array([...]) # labels +1/-1
K_train = qml.kernels.kernel_matrix(X_train, X_train, kernel)
```

This yields an $M_{\text{train}} \times M_{\text{train}}$ matrix. Similarly compute K_{test}

Training SVM with Precomputed Quantum Kernels

With the kernel matrix in hand, we use scikit-learn's SVM with a precomputed kernel. In Python:

```
clf = svm.SVC(kernel='precomputed')  
clf.fit(K_train, Y_train)  
y_pred = clf.predict(K_test)
```

The SVM solver will treat $K_train[i,j]$ as the feature inner products. After fitting, one can retrieve support vectors via **clf.support** if desired. The prediction y_pred will be an array of +1/-1 predictions on the test set.

It is also possible to integrate PennyLane's differentiable capabilities by defining a parameterized kernel and optimizing parameters via gradient descent, but here we keep a fixed feature map.

Example: Kernel SVM in PennyLane

Below is a compact example that ties it together for a toy 2D dataset:

```
import pennylane as qml
from pennylane import numpy as np
from sklearn import svm

# Example dataset
X_train = np.array([[0.1, 0.2], [1.1, -0.4], [0.0, 0.9], [1.0, 0.5]])
Y_train = np.array([+1, +1, -1, -1])
X_test = np.array([[0.2, 0.1], [0.9, 0.0]])

# Define device and feature map
dev = qml.device("default.qubit", wires=2)

@qml.qnode(dev)
def kernel_circuit(x1, x2):
    qml.templates.AngleEmbedding(x1, wires=[0,1])
    qml.adjoint(qml.templates.AngleEmbedding)(x2, wires=[0,1])
    return qml.probs(wires=[0,1])[0]

kernel = lambda a, b: kernel_circuit(a, b)

# Compute kernel matrices
K_train = qml.kernels.kernel_matrix(X_train, X_train, kernel_function=kernel)
K_test = qml.kernels.kernel_matrix(X_test, X_train, kernel_function=kernel)

# Train SVM
```

Discussion of Implementation

In practice, one must consider:

1. Noise and shots: If using real hardware or finite shots, each entry K_{ij} is a noisy estimate. It is common to run many shots (e.g. 1000) to reduce variance. One may also average results or use kernel averaging.
2. Data preprocessing: Classical kernels often require normalized or scaled inputs. Similarly, for quantum kernels, it can help to scale data so that features match the expected range of angles.
3. Computational cost: Computing an $N \times N$ kernel matrix on a quantum device takes $O(N^2)$ circuits. For large N this can be slow. One might use a subset of training data or low-rank approximations.
4. Hyperparameters: The QSVM has hyperparameters (e.g. SVM regularization C) that should be tuned by cross-validation. The choice of feature map is also a “hyperparameter” of sorts.

PennyLane implementations

In summary, implementing QSVM with PennyLane involves four steps:

1. choose a feature map and implement it as a QNode;
2. define a kernel function via that QNode;
3. compute kernel matrices with `qml.kernels`;
4. train/predict using an SVM with the precomputed kernels.

Example Applications (Simple Datasets)

We illustrate QSVM on simple datasets to see how it performs and what it learns. We use the pipeline from Chapter 4. For visualization, we consider 2D synthetic data which can be plotted.

Toy Classification (Circles and Moons)

A common toy problem is concentric circles: two circles of points in the plane. Linearly this is not separable, but a nonlinear kernel (like RBF or a suitable quantum map) can separate them. We generate data:

```
from sklearn.datasets import make_circles
X, y = make_circles(n_samples=100, factor=0.3, noise=0.05)
y = 2*y - 1 # convert to +-1 labels
```

We then train a QSVM with a simple quantum kernel. For visualization, we can plot the points and decision boundary. Typically, a quantum kernel with an entangling map will succeed in separating the circles, similar to an RBF kernel. (If one uses a linear kernel instead, it will fail.)

Example: Using a 2-qubit entangled feature map $U(\mathbf{x}) = R_Y(x_0)\text{CNOT}R_Y(x_1)\text{CNOT}$, we find that the QSVM achieves near-perfect accuracy on this dataset. A contour plot of the decision function shows concentric-level separation.

Iris Dataset

The Iris dataset is a classic benchmark. It has 3 classes of 50 points each, in 4D. We take a two-class subset (e.g. Iris-versicolor vs Iris-virginica) to make a binary problem. We can compare classical SVM vs QSVM:

Classical SVM: with RBF kernel, we typically get 95-100% accuracy on the binary Iris problem. QSVM: using a 2-qubit angle embedding (mapping 4 features to 2 qubits via simple rotations) and RBF on output, one might get similar accuracy.

In practice, many have found that on Iris, a simple RBF classical kernel already saturates performance, so QSVM typically matches it. To illustrate a difference, one can reduce data to 2D and look for small improvements or just use it as an example of applying QSVM to a real dataset.

Dimensionality: QSVM feature maps often require mapping d -dimensional input to some number of qubits n_q . One may use $n_q = \lceil \log_2 d \rceil$ (amplitude encoding) or $n_q = d$ (one feature per qubit). Too many qubits increases circuit depth. Normalization:

Quantum gates assume angles or amplitudes; scaling data into $[0, 2\pi]$ or $[-1, 1]$ may be needed. Train/Test Split: Always evaluate