The Quantum Approximate Optimization Algorithm (QAOA)

From Foundations to Gradient Methods

Morten Hjorth-Jensen

Spring 2026

Outline

1. Combinatorial optimization and the QUBO/Ising mapping
2. Variational quantum algorithms and VQE
3. Definition and structure of QAOA
4. Adiabatic origins and Trotterization
5. MaxCut: explicit construction and gate decomposition
6. Analytic solution at $p = 1$
7. Performance, landscape, and limitations
8. Statistical mechanics and control-theoretic connections
9. Gradient methods for QAOA optimization
10. Summary and outlook

The Optimization Problem

Many problems in physics and computer science reduce to optimizing a cost function over discrete variables:

$$\min_{z \in \{0,1\}^n} C(z)$$

Concrete examples:

- Ground-state energy minimization in spin models
- MaxCut, satisfiability (SAT), traveling salesman
- Scheduling, portfolio optimization

These are in general **NP-hard**: exact solutions scale exponentially with system size n. Approximate algorithms aim for a ratio

$$\alpha = \frac{C(x^*)}{C_{\max}} \lesssim 1.$$

QUBO Formulation

A natural discrete framework is Quadratic Unconstrained Binary Optimization (QUBO):

$$C(x) = x^T Q\, x = \sum_{i,j} Q_{ij}\, x_i x_j, \qquad x^* = \arg \min_{x \in \{0,1\}^n} C(x).$$

Mapping to Ising variables $z_i = 2x_i - 1 \in \{-1, +1\}$ turns the cost function into a *Hamiltonian*:

$$\hat{H}_C = \sum_{i<j} J_{ij}\, Z_i Z_j + \sum_i h_i\, Z_i.$$

This QUBO $\leftrightarrow$ Ising equivalence enables quantum optimization: the ground state of $\hat{H}_C$ encodes the optimal solution. It is also the basis for both quantum annealing and QAOA.

MaxCut: A Benchmark Problem

Given a weighted graph $G = (V, E)$, partition the vertices into two sets to maximize the total weight of cut edges:

$$C(x) = \sum_{(i,j)\in E} w_{ij}\, x_i(1 - x_j).$$

MaxCut is NP-hard and a standard benchmark for quantum optimization.

Quantum Hamiltonian. Using $x_i = (1 - Z_i)/2$:

$$\hat{H}_C = \frac{1}{2} \sum_{(i,j)\in E} w_{ij}(I - Z_i Z_j).$$

The ground state of $\hat{H}_C$ encodes the optimal cut.

Classical baseline. The Goemans–Williamson semidefinite programming algorithm achieves approximation ratio $\alpha \approx 0.878$.

Classical Optimization Methods

Solving the optimization problem classically:

- **Exact methods:** branch-and-bound, integer programming — exponential worst-case cost.
- **Approximation algorithms:** semidefinite programming (Goemans–Williamson, $\alpha \approx 0.878$ for MaxCut).
- **Heuristics:** simulated annealing, genetic algorithms, greedy methods — no guarantee on quality.

Quantum computing offers a potentially different route: use quantum interference and entanglement to explore the configuration space coherently. QAOA is the leading near-term proposal.

Variational Quantum Algorithms (VQAs)

The variational principle underpins a broad class of **hybrid quantum-classical** algorithms:

1. Prepare a parameterized quantum state $|\psi(\boldsymbol{\theta})\rangle = U(\boldsymbol{\theta})|0\rangle$.
2. Measure the expectation value $E(\boldsymbol{\theta}) = \langle\psi(\boldsymbol{\theta})|\hat{H}|\psi(\boldsymbol{\theta})\rangle$ on the quantum device.
3. Update $\boldsymbol{\theta}$ with a classical optimizer.
4. Repeat until convergence.

The quantum device handles the state preparation and measurement; the classical computer handles the optimization.

Variational Quantum Eigensolver (VQE)

VQE is the prototypical VQA for finding ground-state energies:

$$E_0 \approx \min_{\boldsymbol{\theta}} \langle \psi(\boldsymbol{\theta}) | \hat{H} | \psi(\boldsymbol{\theta}) \rangle.$$

Key ingredients:

- A **flexible ansatz** $U(\boldsymbol{\theta})$ — the circuit structure is chosen by the user.
- Measurement of expectation values via Pauli decomposition.
- A classical optimizer (gradient-based or gradient-free).

QAOA is a *special case* of VQE where the ansatz is **problem-inspired** and fixed by the structure of the cost Hamiltonian, rather than chosen freely.

Motivation for QAOA

We want to solve

$$\min_{z\in\{0,1\}^n} C(z)$$

using a quantum circuit that is both *shallow* (NISQ-compatible) and *structured* (informed by the problem).

Two key observations:

- The cost function can be encoded as a diagonal Hamiltonian $H_C = \sum_z C(z)|z\rangle\langle z|$.
- Alternating evolution under H_C and a mixing Hamiltonian H_M can explore the solution space coherently.

QAOA formalizes this into a variational circuit with $2p$ free parameters.

QAOA: Formal Definition

QAOA constructs the parameterized state

$$|\psi_p(\boldsymbol{\gamma}, \boldsymbol{\beta})\rangle = \prod_{k=1}^{p} e^{-i\beta_k H_M} \, e^{-i\gamma_k H_C} \, |+\rangle^{\otimes n},$$

where:

- H_C: cost Hamiltonian encoding the objective function,
- $H_M = \sum_{i=1}^{n} X_i$: mixing Hamiltonian,
- p: circuit depth (number of layers),
- $\boldsymbol{\gamma} = (\gamma_1, \dots, \gamma_p)$, $\boldsymbol{\beta} = (\beta_1, \dots, \beta_p)$: variational parameters.

The objective is to maximize the expectation value:

$$F_p(\boldsymbol{\gamma}, \boldsymbol{\beta}) = \langle \psi_p | \, H_C \, |\psi_p\rangle, \qquad (\boldsymbol{\gamma}^*, \boldsymbol{\beta}^*) = \arg\max \, F_p.$$

Initial State and Hamiltonians

Initial state. The algorithm starts from the equal superposition of all bitstrings:

$$|+\rangle^{\otimes n} = \frac{1}{2^{n/2}} \sum_{z \in \{0,1\}^n} |z\rangle,$$

which is the ground state of $-H_M$ and represents complete ignorance about the solution.

Cost Hamiltonian. The cost function is encoded as a diagonal operator:

$$H_C = \sum_z C(z)\,|z\rangle\langle z|.$$

In Ising form: $H_C = \sum_{i<j} J_{ij} Z_i Z_j + \sum_i h_i Z_i$.

Mixer Hamiltonian. The standard choice $H_M = \sum_i X_i$ generates transitions between computational basis states and ensures exploration of the search space.

QAOA Circuit Structure

Each of the p layers consists of two unitaries:

$$U_C(\gamma_k) = e^{-i\gamma_k H_C}, \qquad U_M(\beta_k) = e^{-i\beta_k H_M}.$$

The full circuit is:

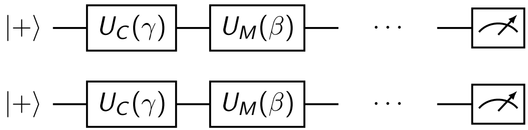

Repeated p times. The approximation ratio is

$$\alpha = \frac{F_p(\gamma^*, \beta^*)}{C_{\max}}.$$

Ansatz Variants

The standard QAOA ansatz can be extended:

- **Multi-angle QAOA:** independent parameters for each gate, increasing expressivity.
- **Adaptive QAOA:** problem-specific ansatz construction, building the circuit layer by layer.
- **Counterdiabatic QAOA:** adds correction terms inspired by counterdiabatic driving to speed up convergence.

All variants share the same hybrid quantum-classical optimization loop and the same underlying physical motivation.

Quantum Adiabatic Algorithm (QAA)

The Quantum Adiabatic Algorithm solves optimization problems by *slow* continuous evolution under a time-dependent Hamiltonian:

$$\hat{H}(s) = (1-s)\,H_M + s\,H_C, \qquad s = t/T \in [0,1].$$

Adiabatic theorem: if the evolution is slow enough (relative to the inverse spectral gap), the system remains in the ground state throughout, so:

$$|\psi(T)\rangle = \mathcal{T}\exp\left(-i\int_0^T \hat{H}(t)\,dt\right)|s\rangle \xrightarrow{T\to\infty} |\text{ground state of } H_C\rangle.$$

The bottleneck is the required time $T \sim 1/\Delta^2$, where Δ is the minimum spectral gap.

Trotterizing the Adiabatic Evolution

Discretize the time interval into p steps of size $\Delta t = T/p$:

$$U(T) \approx \prod_{k=1}^{p} e^{-i\Delta t\, \hat{H}(s_k)}.$$

Apply the Trotter–Suzuki decomposition,
$e^{-i(A+B)\Delta t} \approx e^{-iA\,\Delta t}\, e^{-iB\,\Delta t}$, to each step:

$$U(T) \approx \prod_{k=1}^{p} e^{-i\gamma_k H_C}\, e^{-i\beta_k H_M},$$

where $\gamma_k \sim s_k\,\Delta t$ and $\beta_k \sim (1 - s_k)\,\Delta t$.

$\boxed{\text{QAOA is a Trotterized (discretized) adiabatic evolution.}}$

From Fixed Schedule to Free Parameters

The adiabatic algorithm uses a **fixed** schedule $s(t)$. QAOA replaces it with **free variational parameters**:

$$(\gamma_k, \beta_k) \text{ are optimized classically.}$$

Advantages of the variational approach:

- No need for a slow, hardware-expensive evolution.
- Adapts to hardware noise and connectivity.
- Can outperform the naive adiabatic schedule for finite p.

Limit $p \to \infty$:

$$\lim_{p\to\infty} |\psi_p\rangle \to |\psi_{\text{ground}}\rangle.$$

At finite p we get a shallow, NISQ-compatible circuit.

Connection to Physics

QAOA is equivalent to a **discrete quantum quench sequence**: alternately freeze the system under H_C for time γ_k, then drive it with the mixer H_M for time β_k.

This connects QAOA to:

- **Digital quantum simulation** — Trotterized time evolution.
- **Quantum optimal control** — find the pulse sequence (γ_k, β_k) that maximizes $\langle H_C \rangle$.
- **Many-body dynamics** — the circuit explores entangled spin configurations layer by layer.

MaxCut Cost Hamiltonian (Revisited)

For MaxCut on graph $G = (V, E)$, the classical cost and quantum Hamiltonian are:

$$C(z) = \sum_{(i,j)\in E} \frac{1 - z_i z_j}{2}, \qquad H_C = \sum_{(i,j)\in E} \frac{1 - Z_i Z_j}{2}.$$

This is equivalent to an Ising model:

$$H = -\sum_{(i,j)\in E} J_{ij}\, Z_i Z_j,$$

so QAOA on MaxCut is exactly a quantum exploration of spin configurations on the graph. The ground state encodes the maximum cut.

Gate Decomposition: Cost Unitary

For one edge (i,j), the cost unitary is $e^{-i\gamma(1-Z_iZ_j)/2}$. The non-trivial part is

$$e^{-i\gamma Z_i Z_j},$$

implemented by the standard CNOT–R_Z–CNOT circuit:

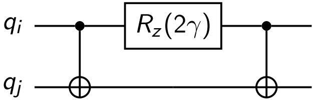

One such two-qubit block is applied for every edge in the graph.

Gate Decomposition: Mixer Unitary

The mixer $e^{-i\beta H_M} = \prod_i e^{-i\beta X_i}$ is a product of single-qubit x-rotations:

$$e^{-i\beta X_i} = R_x(2\beta).$$

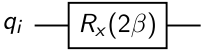

Applied independently to every qubit. Together with the cost unitary, these two building blocks make up a single QAOA layer.

Classical Optimization Loop

Once the circuit is built, QAOA runs as a hybrid loop:

1. **Initialize** parameters (γ, β).
2. **Run** the quantum circuit and measure $\langle H_C \rangle$.
3. **Update** parameters using a classical optimizer (gradient-based or gradient-free).
4. **Repeat** until convergence.

The quantum device is queried repeatedly; the classical computer drives the parameter update. The quality of the solution depends critically on the optimizer choice and initialization.

Setup: QAOA at $p = 1$

Consider two qubits connected by a single edge, so

$$H_C = \frac{1 - Z_1 Z_2}{2}.$$

The $p = 1$ QAOA state is

$$|\psi(\gamma, \beta)\rangle = e^{-i\beta H_M} \, e^{-i\gamma H_C} \, |+\rangle^{\otimes 2}.$$

Apply $U_C(\gamma)$ first: since $Z_1 Z_2$ has eigenvalues ± 1,

$$U_C = e^{-i\gamma/2} \, e^{i\gamma Z_1 Z_2/2}.$$

Then apply $U_M(\beta) = e^{-i\beta(X_1 + X_2)} = R_x(2\beta)^{\otimes 2}$.

Key Computation: Edge Expectation Value

Carrying out the algebra (a standard result), one finds for a single edge:

$$\langle Z_1 Z_2 \rangle = \sin(4\beta)\sin(2\gamma).$$

Therefore the expectation value of the cost Hamiltonian is

$$F(\gamma, \beta) = \langle H_C \rangle = \frac{1}{2}\big(1 - \sin(4\beta)\sin(2\gamma)\big).$$

For a general graph, by linearity:

$$F(\gamma, \beta) = \sum_{(i,j)\in E} \frac{w_{ij}}{2}\big(1 - \sin(4\beta)\sin(2\gamma)\big),$$

where contributions from different edges decouple at $p = 1$.

Optimal Parameters and Approximation Ratio

The expectation value is maximized at

$$\beta^* = \frac{\pi}{8}, \qquad \gamma^* = \frac{\pi}{4},$$

giving

$$\sin(4\beta^*)\sin(2\gamma^*) = 1.$$

For triangle-free regular graphs this yields an approximation ratio

$$\alpha \approx 0.6924,$$

matching known analytic results from Farhi, Goldstone, and Gutmann.

The $p = 1$ circuit captures only *short-range* correlations (depth-1 light cone). Higher p builds long-range correlations and approaches the adiabatic limit.

Locality and General Graphs

An important structural fact: at depth p, the QAOA circuit has a light-cone of radius p.

$$\text{At depth } p: \text{ correlations extend only to distance } p.$$

Consequently:

- For small p, the expectation value of each edge term depends only on its local neighborhood of radius p in the graph.
- This makes analytic computations tractable for small p.
- It also limits performance on problems with long-range correlations unless p is large.

Performance as a Function of Depth

- $p = 1$: shallow circuit, limited accuracy, analytically tractable.
- Larger p: better approximation ratio, closer to adiabatic limit.
- $p \to \infty$: exact solution recovered.

Trade-off:

$$\text{accuracy} \longleftrightarrow \text{circuit depth (hardware cost)}.$$

On NISQ devices, gate errors and decoherence limit the practical depth. Noise mitigation and hardware-aware ansätze are active research topics.

Energy Landscape

The objective function

$$E(\gamma, \beta) = \langle \psi(\gamma, \beta) | H_C | \psi(\gamma, \beta) \rangle$$

is a smooth, 2π-periodic function of $2p$ parameters.
Properties:

- **Small** p**:** relatively smooth landscape, easy to optimize.
- **Large** p**:** exponentially many local minima; landscape resembles spin-glass free-energy surfaces.

Parameter optimization strategies:

- Good initialization (e.g., from a lower-p solution).
- Gradient-based methods (parameter-shift rule).
- Exploiting symmetries of the cost Hamiltonian.

Barren Plateaus

A central challenge for large parameterized circuits is the **barren plateau** phenomenon: gradients vanish exponentially with system size,

$$\mathrm{Var}[\partial_{\theta_k} E(\boldsymbol{\theta})] \sim \mathcal{O}(b^{-n}), \qquad b > 1.$$

Consequences:

- Training becomes exponentially slow for deep, expressive circuits.
- Gradient-based optimizers stall.

Why QAOA is more resistant:

- The ansatz is **problem-inspired** — the structure of H_C concentrates relevant correlations.
- **Locality** of the circuit (depth-p light cone) limits the spread of barren-plateau effects.

Optimization as Physics

Minimizing $C(z)$ is equivalent to finding the ground state of

$$H_C = \sum_z C(z) \, |z\rangle\langle z|,$$

which is a classical Ising Hamiltonian. This maps:

- combinatorial optimization $\leftrightarrow$ spin-system ground state,
- QAOA $\leftrightarrow$ controlled quantum exploration of spin configurations.

Physical picture:

- H_C encodes the energy landscape.
- H_M induces quantum tunneling between spin configurations.
- The QAOA layers drive the system toward low-energy states via constructive quantum interference.

Real-Time vs Imaginary-Time Evolution

Classical statistical mechanics uses the **Gibbs state**:

$$\rho_{\mathrm{Gibbs}} \propto e^{-\beta H}, \qquad Z = \mathrm{Tr}(e^{-\beta H}).$$

QAOA instead uses **real-time** unitary evolution:

$$|\psi\rangle = e^{-i\beta H_M} \, e^{-i\gamma H_C} \, |+\rangle.$$

The correspondence $t \to -i\beta$ links:

- imaginary time $\to$ Gibbs sampling / simulated annealing,
- real time $\to$ quantum interference / coherent dynamics.

QAOA exploits quantum *interference* rather than thermal fluctuations, which is the fundamental distinction from simulated annealing.

Relation to Annealing Algorithms

Method	Mechanism	Evolution
Simulated annealing	thermal fluctuations	stochastic
Quantum annealing	tunneling	continuous unitary
QAOA	coherent interference	discrete unitary

QAOA uses *controlled* coherent evolution with variational parameters, making it more flexible than fixed-schedule annealing while remaining tractable on finite-depth quantum circuits.

Control-Theoretic Interpretation

From a quantum control perspective, QAOA implements a piecewise-constant control protocol:

$$U(\boldsymbol{\theta}) = \mathcal{T} \exp\left(-i \int_0^T \left[u_1(t)\, H_C + u_2(t)\, H_M\right] dt\right),$$

where u_1, u_2 are the piecewise-constant controls (γ_k, β_k). The optimization objective

$$\max_{\boldsymbol{\gamma},\boldsymbol{\beta}} \langle \psi(T) | H_C | \psi(T) \rangle$$

is a **quantum optimal control problem**.

As $p \to \infty$, QAOA converges to a continuous Schrödinger evolution and the control problem approaches the adiabatic limit.

The Gradient Problem

Optimization of the $2p$ parameters requires the gradient

$$\nabla E = \left(\frac{\partial E}{\partial \gamma_1}, \frac{\partial E}{\partial \beta_1}, \ldots, \frac{\partial E}{\partial \gamma_p}, \frac{\partial E}{\partial \beta_p} \right).$$

If the cost Hamiltonian has a Pauli decomposition $H_C = \sum_{\alpha=1}^{K} h_\alpha P_\alpha$, then

$$E(\boldsymbol{\theta}) = \sum_{\alpha=1}^{K} h_\alpha \langle \psi(\boldsymbol{\theta}) | P_\alpha | \psi(\boldsymbol{\theta}) \rangle,$$

and the gradient becomes a sum over Pauli-string derivatives.

Central question: How does one evaluate $\partial_{\theta_k} \langle P_\alpha \rangle$ efficiently on hardware?

General Derivative Structure

Let $|\psi_0\rangle = |+\rangle^{\otimes n}$ and $U(\boldsymbol{\theta})$ be the full QAOA unitary. Then

$$E(\boldsymbol{\theta}) = \langle\psi_0|U^\dagger(\boldsymbol{\theta})\, H_C\, U(\boldsymbol{\theta})|\psi_0\rangle.$$

Each parameter θ_k appears in a gate $U_k(\theta_k) = e^{-i\theta_k B_k}$ with generator $B_k \in \{H_C, H_M\}$. Writing $U = U_{\text{after}}\, U_k(\theta_k)\, U_{\text{before}}$, all gradient methods reduce to evaluating derivatives of $\langle U^\dagger B_k U\rangle$.

Why is this nontrivial?

- Hardware cannot apply $H|\psi\rangle$ directly as a gate.
- Each gradient evaluation costs $\mathcal{O}(p)$ circuit calls.
- Many-term Hamiltonians ($K \gg p$) can dominate cost.

A Concrete Test Case: XXZ Spin Chain

The five gradient methods are compared on the XXZ Hamiltonian:

$$H_{\mathrm{XXZ}} = \sum_{i=1}^{L-1} \left[J_{xx} (\sigma_i^x \sigma_{i+1}^x + \sigma_i^y \sigma_{i+1}^y) + J_z \, \sigma_i^z \sigma_{i+1}^z \right] + \sum_{i=1}^{L} h_z \, \sigma_i^z ,$$

with $J_{xx} = 1.0$, $J_z = 1.5$, $h_z = 0.2$, and $L = 6$, $p = 2$, giving $N = 2p = 4$ variational parameters and approximately $K \approx 14$ Pauli terms.

The coupling ratio $h_{\max}/h_{\min} = 7.5$ makes this a non-trivial benchmark for methods that exploit Hamiltonian structure.

Method 1: Exact Gradient

The exact gradient formula follows directly from differentiating the unitary evolution:

$$\frac{\partial E}{\partial \theta_k} = i\langle\psi|H_C\, U^\dagger B_k U|\psi\rangle - i\langle\psi|U^\dagger B_k U\, H_C|\psi\rangle.$$

- **Implementation:** apply the Hamiltonian MPO to intermediate states and compute overlaps.
- **Cost:** N circuit evaluations $+$ N Hamiltonian applications.

\+ Exact gradient, minimal approximation error.

– Requires $H|\psi\rangle$ as a primitive operation — *not directly implementable on standard quantum hardware*.

Method 2: Standard Parameter-Shift Rule

The standard parameter-shift rule evaluates $\partial E / \partial \theta_k$ through two shifted circuit evaluations:

$$\frac{\partial E}{\partial \theta_k} = c[E(\theta_k + s) - E(\theta_k - s)]\,,$$

for appropriate shift s and prefactor c depending on the spectrum of the generator B_k.

+ Hardware compatible — requires only expectation-value measurements of shifted circuits.
+ Simple, widely used baseline.
− Treats the generator uniformly; does not exploit Hamiltonian structure.
− Can be inaccurate when the Pauli decomposition is highly non-uniform.

Method 3: Per-Pauli Parameter Shifts

Method 3 decomposes the generator $B = \sum_\alpha h_\alpha P_\alpha$ and applies parameter shifts at the level of *individual Pauli terms*, adapting the shift to each coefficient h_α.

+ Improved accuracy for Hamiltonians with non-uniform coefficients.
+ Finer control over gradient estimation.
– More circuit evaluations than Method 2 (one pair per Pauli term).
– More bookkeeping; higher implementation complexity.

Best suited when $h_{\max}/h_{\min}$ is large and accuracy matters more than circuit count.

Method 4: Commuting-Group Optimization

Method 4 improves on Method 3 by grouping Pauli terms that *commute* with each other. Commuting terms can share a measurement basis or be shifted collectively, reducing redundant evaluations.

Benefits:

- Reduced measurement overhead.
- Fewer shifted circuit evaluations.
- Retains much of Method 3's accuracy gain.

In practice, this reduces cost by roughly **60%–80%** compared to a naïve per-Pauli strategy, making it a practical route to Hamiltonian-aware gradients.

Method 5: Adaptive Hybrid Strategy

Method 5 combines ideas from all previous methods:

- Use exact or per-Pauli shifts where Hamiltonian structure matters most (large $|h_\alpha|$).
- Use cheaper uniform shifts where individual term effects are small.
- Balance accuracy and circuit cost *adaptively*.

It is the practical middle ground:

Method 1	exact, hardware-incompatible
Method 2	simple, hardware-compatible, least accurate
Method 3	accurate, expensive
Method 4	accurate, cost-reduced
Method 5	**adaptive compromise**

Comparison of Gradient Methods

For the XXZ chain example ($L = 6$, $p = 2$, $K \approx 14$):

- Non-uniform couplings ($h_{\text{max}}/h_{\text{min}} = 7.5$) make Methods 3–5 **more accurate** than Method 2.
- Commuting-group and adaptive strategies reduce cost by **60%–80%**.
- Method 2 remains the **easiest standard implementation**.
- Methods 3–5 become attractive when Hamiltonian coefficients vary strongly.

Key insight: there is no single best method for all Hamiltonians. The choice depends on the coupling structure, available circuit budget, and required accuracy.

Possible Extensions for Gradient Methods

Active research directions:

- Explicit derivation of parameter-shift rules for QAOA generators beyond two-eigenvalue spectra.
- Measurement grouping strategies for large Pauli decompositions.
- Shot-noise analysis for each method.
- Tensor-network (MPO) simulation costs for classical benchmarking.
- Application to larger spin models and combinatorial Hamiltonians beyond MaxCut.

Summary

QAOA in one picture:

$$\underbrace{|+\rangle^{\otimes n}}_{\text{initial state}} \xrightarrow{\prod_k e^{-i\gamma_k H_C} e^{-i\beta_k H_M}} |\psi_p\rangle \xrightarrow{\text{measure}} \langle H_C\rangle \xrightarrow{\text{optimize}} (\boldsymbol{\gamma}^*, \boldsymbol{\beta}^*)$$

Key messages:

- QAOA encodes combinatorial optimization as Ising Hamiltonians.
- It is a Trotterized adiabatic evolution with variational parameters.
- Analytic results are available at small p; the $p = 1$ MaxCut solution gives $\alpha \approx 0.693$.
- The algorithm connects quantum computing with statistical physics and quantum optimal control.
- Gradient computation is a practical bottleneck; per-Pauli and adaptive methods can outperform the standard parameter-shift rule for structured Hamiltonians.

Outlook and Open Questions

- **Quantum advantage:** Does QAOA outperform the best classical algorithms (e.g., Goemans–Williamson) for any problem? Open for $p \geq 2$.
- **Scalability:** Noise and barren plateaus limit practical depth on NISQ devices; fault-tolerant hardware may unlock large p.
- **Beyond MaxCut:** Application to constrained optimization, machine learning, and many-body physics problems.
- **Better mixers:** Problem-specific mixers that preserve feasibility constraints outperform the standard $\sum X_i$.
- **Classical simulation:** For which problem classes can QAOA be efficiently simulated classically?

QAOA remains one of the most actively studied algorithms in quantum computing, at the intersection of optimization, physics, and information theory.

Thank you. Questions?